

Deep Q Learning: From Paper to Code

Coding The Deep Q Network

What We Need to Accomplish

- 3 Convolutional layers, 2 fully connected layers
 - C1: 32, 8x8, stride 4; C2: 64, 4x4, stride 2, C3: 64, 3x3, stride 1; FC: 512
- Will need function to find input size to FC layer
- RMSProp optimizer, MSE Loss
- Should also support model checkpointing