

Deep Q Learning: From Paper to Code

Temporal Difference Learning

Last Time ...

- Explore for long term or exploit for short term?
- Many solutions \rightarrow epsilon greedy
- Occasionally random, occasionally greedy

Refining Value Function Estimate

$$v_{\pi}(s) = \sum_a \pi(a, s) \sum_{s', r} p(s', r | s, a) [r + \gamma v_{\pi}(s')]$$

Estimate of value of a policy



Refine estimate


New estimate = old estimate + step size (target – old estimate)

Target?

New estimate = old estimate + step size (target – old estimate)



Direction of update



May be noisy




R_t

Step Size?

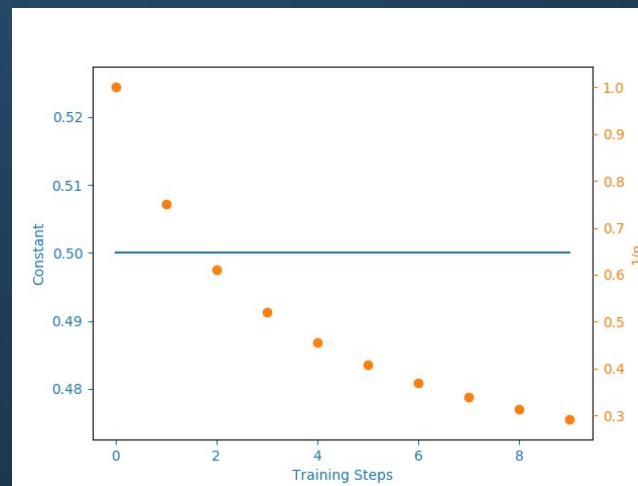
New estimate = old estimate + step size (target – old estimate)



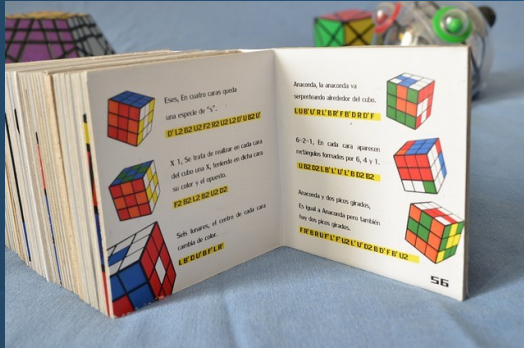
Control rate of change



Many forms



Update Frequency



Algorithm dependent




Monte Carlo → end of episode

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Sum of discounted rewards that follow current time

New estimate = old estimate + step size (target – old estimate)

Temporal Difference Learning




Q learning



Update at each time step




Online learning



Non episodic tasks

Temporal Difference Updates

$$V(s_t) = V(s_t) + \alpha(R_{t+1} + \gamma V(s_{t+1}) - V(s_t))$$




Using earlier estimate to update

Bootstrapping

Importance of initial estimate?

On Convergence



Fixed Π and small $\alpha \rightarrow$ converge with good coverage

Hundreds of visits or so



How to use in Q learning?

Application to Q Learning

Have to consider action-value function!

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha(R_{t+1} + \gamma \max_a Q(s_{t+1}, a_{max}) - Q(s_t, a_t))$$

Compare to

$$V(s_t) = V(s_t) + \alpha(R_{t+1} + \gamma V(s_{t+1}) - V(s_t))$$

Use maximal Q for update equation

In the Frozen Lake



Small number of states, actions

	Action 1	Action 2	Action n
State 1	$Q(1,1)$	$Q(1,2)$	$Q(1,n)$
State 2	$Q(2,1)$	$Q(2,2)$	$Q(2,n)$
State m	$Q(m,1)$	$Q(m,2)$	$Q(m,n)$

Tabular Learning Method

Implementation Details



Epsilon-greedy action selection


For a given state (s) look up Q for each action (a)

Max action or random action

Use reward to update Q for s, a

Dont' reset table \rightarrow coverage of all states

Off Policy Learning



Epsilon greedy to update greedy

Off Policy Learning

SARSA \rightarrow On policy

Check out my YouTube for more

Q Learning Algorithm

Initialize Q for all states s and actions a

Initialize $\alpha=0.001$ $\gamma=0.9$ $\epsilon_{max}=1.0, \epsilon_{min}=0.01$

Repeat for n_episodes

Initialize state s

For each step of episode

Choose a with epsilon-greedy

Perform a, get new state s' and reward r

$$Q(s, a) = Q(s, a) + \alpha(r + \gamma \max_a Q(s', a) - Q(s, a))$$

$s = s'$

Agent should be a class

Q is a dictionary

Decrement epsilon over time

Separate files

Plot average score over 100 games

500,000 total games



Summary

- Temporal difference learning is online
- Bootstrap learning
- Q learning is tabular and off policy
- Achieve 70% win rate, up from 20%

Up Next

